


LDSCI5247 Foundations of Data Science AE2

Assessment Details

Course Leader:	Anees Baqir
First, Second, or Third Sitting:	First
Issue Date:	8 September 2025
Assessment Type:	Code & Report
Assessment Title:	Classifying Countries Based on Outbound Student Migration Trends
Restrictions on Time/Length:	1500 Words
Assessment Weighting:	50%
Hand-in Deadline:	5 December 2025
Planned Feedback Deadline:	28 calendar days after hand-in deadline or last presentation date
File Format Accepted:	.html export of Python Notebook
Mode of Submission:	Online (Canvas)
Anonymous Submission:	Yes
Marking Scheme:	Categorical

	Prohibited	Use of AI is prohibited on this assignment.
-------------------------------------------------------------------------------------	-------------------	---------------------------------------------

Assessment Task

Introduction

International student migration reflects not only educational preferences but also broader socio-economic and policy trends across countries. Some countries consistently exhibit high outbound student mobility, while others show relatively lower rates. In this assessment, you will classify countries based on whether they are **High-Senders** or **Low-Senders** of students, based on engineered features derived from migration trends between 2019 and 2023.

You will develop and apply classification models to predict country types using relevant statistical and temporal features. This task allows you to demonstrate your ability to

perform data integration, feature engineering, binary classification, and model interpretation.

Question 1 Feature Engineering and Labelling (20 marks)

Begin by loading the student migration dataset into a Pandas DataFrame. Clean the dataset to remove or impute any missing or inconsistent values, particularly for origin countries and outbound counts. Focus on the **2019–2023** period.

Your first goal is to engineer a new dataset where each **origin country** is labelled as either a **High-Sender** (1) or **Low-Sender** (0) of international students. Use a rule-based threshold to assign this label: for instance, compute the **average annual outbound student count** per country over the 5-year period, and label countries above the **75th percentile** as *High-Senders* and the rest as *Low-Senders*.

Then, extract and compute time-series based features that may help in predicting this classification. Examples of useful features include:

- Linear trend (slope) of outbound students from 2019 to 2023
- Year-over-year growth rate
- Standard deviation or volatility in outbound numbers
- Max or min student flow in the period
- Average annual flow

Once these features are generated, ensure they are combined into a clean tabular format where each row corresponds to one origin country, and no missing values remain. This dataset will serve as the input for your classification models.

Question 2 Model Training and Classification (35 marks)

Using the labelled dataset created in Task 1, build classification models to predict whether a country is a **High-Sender** or **Low-Sender**. Use the following four algorithms for this task:

- Linear Probability Model
- Logistic Regression
- k-Nearest Neighbors (kNN)
- Decision Tree Classifier

Split your data into a training set and testing set (e.g., 70/30 split). Train each model using the training data and evaluate performance using **accuracy**, **precision**, **recall**, and **F1-score** on the test set. Use a classification report and other appropriate visual tools to summarise model results. After training, compare the performance of all four models and explain which performed best in terms of predictive accuracy and balance between false positives and false negatives.

Question 3 Model Interpretation and Insights (35 marks)

In this final task, reflect on the results of your models. Which features were most predictive of a country being a High-Sender? Use appropriate model interpretation tools, such as feature importance in decision trees or logistic regression coefficients to support your discussion.

Were there any surprising patterns in which countries were classified as High-Senders or Low-Senders? Consider potential factors such as economic development, language, regional proximity, or cultural preferences that might influence outbound mobility and are not directly captured in the dataset.

Lastly, assess how well your models generalise. Do any models show signs of overfitting? Are there alternative features you might include (e.g., external data such as GDP, population, or education index) to improve predictions in future iterations?

Your discussion should demonstrate a critical understanding of both the modelling process and the data domain.

Note: 10 marks are awarded for the English proficiency of your report.

The dataset required for this assessment will be provided by the instructor and made available exclusively through the course Canvas page. Please do not download external versions or alternative sources of the dataset.

Assessment Criteria

The assessment should be completed in a Python notebook using Markdown cells for the written components of the task. You must remove any unnecessary code or markdown cells from the notebook before submitting the assignment.

Assessment Criteria

70 or higher	There was evidence of the ability to perform all tasks correctly. The demonstration of the methods was excellent, coherent, well documented, and clearly explained.
60-69	There was evidence of ability to perform some tasks correctly. The demonstration of the methods is good, coherent, and reasonably detailed and explained.
50-59	There was evidence of ability to perform some tasks correctly, but the demonstration of the methods was limited, incoherent, not adequately documented and vaguely explained.

- | | |
|------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 40-49 | There was limited evidence of ability to perform tasks. The demonstration of the methods involved significant omissions and produced substantial inaccuracies. |
| 39 or less | Failure to solve the tasks in the assignment. Methods were completely incorrect or absent. |

Submitting Assessments

For this assignment:

- Prepare a single Jupyter Notebook file that includes all the code, text, printed outputs, and data visualisations. Alternatively, prepare a Jupyter Notebook file that contains the code and prepare a .html file with the text and data visualisations.
- The notebook should be developed using Python 3 within a Jupyter environment, with well-commented code and appropriately structured text cells resembling a blog format.
- Do not include any identifying information such as your name.
- Remember to restart the kernel, thoroughly check the notebook for errors, and provide a working link to the dataset used in the analysis.
- Ensure that all code runs smoothly and create a compelling narrative by integrating analysis, visualisations, and reflections.
- Your submission may include auxiliary Python files (*.py), where you packaged, e.g., some common functions and imported them in your notebook. In this case, submit a .zip file with all documents included.
- Do not include any dataset files in your submission.

If you encounter issues with submission:

- Check the assessment details table at the top of the assessment brief to be sure you are submitting a permitted file format. Avoid zip files (unless explicitly required) and scanned PDFs. Use the 'add files' function to submit multiple files instead of a zip file.
- Make sure you have ticked the agreement box at the bottom of your Canvas submission page (scroll down if you don't see it). This will enable you to select 'Submit Assessment.'
- Try changing web browsers.

If you still cannot submit, e-mail a copy of your assignment before the deadline to student.assessments@nulondon.ac.uk along with screenshots of the problem on Canvas, showing a timestamp.

If your assessment requires anonymous submission (see the assessment details table at the top of the assessment brief), be sure you have left your name off of your submission and out of the submission file name.

Please review the submitted file to ensure that everything is in order. To turn on notifications for submission confirmation emails in your Canvas settings: Account > Notifications > Turn on the bell for 'All submissions.' In the app this is via Settings > Email Notifications > All submissions.

Marking

The University uses two assessment marking schemes – one for undergraduate and one for postgraduate – to mark all taught programmes leading to an award of the University.

More detailed information on the assessment marking scheme and the criteria can be found in the Course Syllabus, available on the University's VLE.

Learning Outcomes

This assessment will enable students to demonstrate in full or in part the learning outcomes identified in the Course Descriptor.

On successful completion of this assessment, students should be able to:

Knowledge and Understanding

K1b	Demonstrate knowledge and critical of well-established data representation and transformation concepts for multi-dimensional data.
K2b	Demonstrate ability to identify feasible operations and transformation on data, and their relationships in a data processing pipeline.
K3b	Demonstrate knowledge and critical understanding of plotting and visualising data.

Subject-Specific Skills

S2b	Apply the data science theory learnt in class (e.g., well-established data transformation techniques) in an appropriate manner to a given dataset.
S3b	Identify the correct choice of appropriate data transformation techniques.

Transferable Skills

T2b	Identify, transform, evaluate, and plot accordingly from a dataset.
-----	---------------------------------------------------------------------

T3b	Demonstrate a sound technical proficiency in written English and skill in selecting vocabulary so as to communicate effectively to specialist and non-specialist audiences.
-----	-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Accessing Feedback

Students can expect to receive feedback on all summative coursework within 28 calendar days of the submission deadline or, if applicable, the last oral assessment date, whichever later. The 28 calendar day deadline does not apply to work submitted late. Feedback can be accessed through the assessment link on the Canvas course page.

Late Submissions

Please ensure that you submit your assignment well before the deadline to avoid any late penalties, as a submission made exactly on the deadline will be considered late. Please keep in mind that there may be differences between your computer's clock and the server time, which can cause discrepancies, and that Canvas may take some time to process your submission.

Your Canvas submission portal displays two due dates: one is the deadline for your assignment, and the second is the latest possible date by which your assignment can be submitted late. Please make sure you submit by the assessment deadline in order to avoid late penalties.

If assessments are submitted late without approved Extenuating Circumstances, there are penalties:

- For assessment elements submitted up to one day late, any passing mark will receive 10 marks deducted or a threshold pass (40% for undergraduate students, 50% for postgraduate students), whichever is higher. Any mark below 40% for undergraduate students and below 50% for postgraduate students will stand.
- Students who do not submit their assessment within one day of the deadline, and have no approved Extenuating Circumstances, are deemed not to have submitted and to have failed that assessment element. The mark recorded will be 0%.
- For assessment subelements, late submission will result in non-submission penalties deducted according to the marking criteria above.

For further information, please refer to [AQF7 Part C in the Academic Handbook](#).

Extenuating Circumstances

The University's Extenuating Circumstances (ECs) procedure is in place if there are genuine circumstances that may prevent a student from submitting an assessment. If the EC application is successful, there will be no academic penalty for missing the published submission deadline.

Students are normally expected to apply for ECs in advance of the assessment deadline. Students may apply for consideration of ECs retrospectively if they can provide evidence that they could not have done so in advance of the deadline. All applications for ECs must be supported by independent evidence.

Successful EC applications for live oral assessments, including vivas, will result in a deferral of the oral to be organized by faculty, students, and Timetabling for a date as close as possible to the original presentation date. The deadline for supplementary materials, if assigned, will be carried forward by the length of the oral assessment extension.

Missing an oral assessment, including a compulsory viva, without an approved EC will result in a non-submission for the entire assessment and, accordingly, a recorded mark of 0%.

Students are reminded that the ECs procedure covers only short-term issues (within 21 days leading to the submission deadline) and that if they experience longer-term matters that impact on learning then they must contact [Student Support](#) for advice.

Under the Extenuating Circumstances Policy, students may defer an assessed element on only one occasion and may request an extension on a maximum of two occasions.

For further information, please refer to the [Extenuating Circumstances Policy](#) in the Academic Handbook.

Academic Misconduct

You must submit work for assessment purposes that is your own and meets good academic practice. Assessments must be completed strictly in accordance with the instructions outlined in the assessment brief. This includes ensuring that your work is appropriately referenced.

It is important to understand if artificial intelligence (AI) is permitted to be used or not, and if yes, the conditions for its use. For clarity, the assessment brief expressly states if AI can be used, and the parameters of its permitted usage. You are not permitted to use AI (in any form) if the assessment brief does not expressly state that usage is permitted.

Any concerns with the academic integrity of a submission will be addressed in accordance with the [Academic Misconduct Policy](#). This policy gives details of the different types of misconduct and the steps the University may take when a concern arises. Action may result in academic penalties being applied directly to you. You are advised to review this policy in full before completing and submitting any work for marking purposes. The University may also take steps to ascertain the authenticity of a submitted piece of work in cases of relevant concern, such as by holding a viva.

Version History

Title: Assessment Brief Template					
Approved by: The Quality Team					
Version number	Date approved	Date published	Owner	Location	Proposed next review date
4.0	March 2023	March 2023	Registrar	VLE/ Faculty Resources Page	March 2024
3.0	August 2022	August 2022	Registrar	VLE, Faculty Resources Page	July 2023
2.3	December 2021	December 2021	Registrar	VLE	August 2022
2.2	August 2021	August 2021	Registrar	VLE	August 2022
2.1	September 2020	September 2020	Registrar	VLE	August 2021
2.0	September 2020	September 2020	Registrar	VLE	August 2021
1.0	August 2019	August 2019	Registrar	VLE	August 2020
Referenced documents	AQF7 Academic Regulations for Taught Awards; Extenuating Circumstances Policy; Academic Misconduct Policy; Course Syllabus				
External Reference Point(s)	UK Quality Code Theme: Assessment				