

# Week 7: Regression 1

## Univariate Regression

Dr Giuseppe Brandi

Northeastern University London

# What is regression?

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

A type of supervised learning, where each input example has a known target:

- If target is **discrete**, then it is a classification problem
  - "Which one?"
- If target is **continuous**, then it is regression
  - "How many?" or "How much?"

# Example problems

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

“Predict

- CO<sub>2</sub> emissions<sup>1</sup>
- Bicycle traffic

in London given *datetime*, *weather season*, etc.”

So, we collect samples – **both** input and target values – and we wish to understand a trend.



---

<sup>1</sup>Note that targets are continuous random variables.

# Linear regression, Part I

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

The simple **linear regression** model involves a linear combination of input variables:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x_1 + w_2x_2 + \cdots + w_dx_d$$

where  $\mathbf{x} = (x_1, \dots, x_d)^T$  is a  $d$ -dimensional input vector. The key idea is that  $y(\mathbf{x}, \mathbf{w})$  is a **linear function** of parameters  $\mathbf{w}$

# Fitting a straight line

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

The simplest example is to fit a straight line to data:

$$y = w_0 + w_1 x$$

Parameter  $w_0$  is known as the **intercept** of the line;  
and  $w_1$  as the **slope** of the line

Let's try an example...

# Fitting a straight line

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

## Generating sample data

```
1 import numpy as np
2
3 # Create a random number generator (with seed)
4 rng = np.random.RandomState(123456789)
5 # Generate 50 random values for x, scaled x10
6 x = 10 * rng.rand(50)
7 # Compute y(x), adding random noise
8 y = 2 * x - 5 + rng.randn(50)
```




# Fitting a straight line

Consider the task of approximating the line  $y(x) = -5 + 2x$ .

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# Solving it in Python

## Fit a straight line

```
1 from sklearn.linear_model import LinearRegression
2
3 # Create a linear regression model
4 model = LinearRegression(fit_intercept=True)
5 # Fit the training data
6 model.fit(x[:, np.newaxis], y)
7 # Predict y-value for 1,000 data points
8 u = np.linspace(0, 10, 1000)
9 v = model.predict(u[:, np.newaxis])
```




# Fitting a straight line

Consider the task of approximating the line  $y(x) = -5 + 2x$ .

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# Solving it in Python

## Fit a straight line

```
1 print('y(x) = {:.3f} + {:.3f}x'.format(  
2     model.intercept_,  
3     model.coef_[0]))
```

y(x) = -5.259 + 2.041x

The result is pretty close. *Is it only about straight lines?*

# Linear regression, Part II

The simple linear regression model is also a linear function of input variables  $x$ , which is a limiting factor.

So, we consider linear combinations of non-linear functions:

$$y(x, w) = w_0 + w_1 \phi_1(x) + w_2 \phi_2(x) + \cdots + w_{m-1} \phi_{m-1}(x)$$

where  $\phi_i(x)$  are known as **basis functions**. The size of the model (that is, the number of parameters) is  $m$ .

# Linear regression, Part II

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

For example, if  $\phi_j(x) = x^j$  then:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \cdots + w_{m-1}x^{m-1}$$

is a **polynomial regression** model.

In a similar fashion, we can use a Gaussian basis function  
(see our *handbook*)

# Polynomial curve fitting

Let's construct a more complicated example, fitting a **curve**, using the function  $y(x) = \sin(2\pi x)$ .

We will generate 10 sample values for  $x$  and  $y(x)$ , adding random normal noise to the latter.

## Generating sample data

```
1 # Generate 10 equally-spaced values in range 0 to 1
2 x = np.linspace(0, 1, 10)
3 # Compute y(x), adding noise from
4 # a normal distribution N(0, 0.3)
5 y = np.sin(2 * np.pi * x) +
6     rng.normal(scale=0.3, size=10)
```




# Polynomial curve fitting

Let's construct a more complicated example, fitting a **curve**, using the function  $y(x) = \sin(2\pi x)$ .

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# Polynomial curve fitting

Let's construct a more complicated example, fitting a **curve**, using the function  $y(x) = \sin(2\pi x)$ .

Our task is to discover the function  $y(x)$  from a finite dataset by **fitting the training data** with a polynomial function.

Let's consider, for now, a polynomial of *degree* (or *order*) 4:

$$y = w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

The polynomial will be determined by  $w_1, \dots, w_4$ .

# Polynomial curve fitting

Let's construct a more complicated example, fitting a **curve**, using the function  $y(x) = \sin(2\pi x)$ .

## Solving it in Python


```
1 from sklearn.preprocessing import PolynomialFeatures
2 from sklearn.pipeline import make_pipeline
3
4 # A polynomial regression model of degree 4
5 model = make_pipeline(PolynomialFeatures(4),
6                         LinearRegression())
7 model.fit(x[:, np.newaxis], y)
8 # Predict y-value for 1,000 data points
9 u = np.linspace(0, 1, 1000)
10 v = model.predict(u[:, np.newaxis])
```

# Polynomial curve fitting

Let's construct a more complicated example, fitting a **curve**, using the function  $y(x) = \sin(2\pi x)$ .

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# How does it work?

The values of the coefficients (in our case,  $w_1, \dots, w_4$ ) are determined by minimising an **error function**.

The error measures the misfit between  $\hat{y}$  and  $y$ , the expected result.

A widely used error function is the sum-of-squares error:

$$E(w) = \frac{1}{2} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$


where  $n$  is the number of input samples.


# How does it work?

The error measures the misfit between  $\hat{y}$  and  $y$ , the expected result.

Week 7:  
Regression 1


Dr Giuseppe  
Brandi





# What has happened when $m = 10$ ?


We have 10 sample data points (the blue points in the graph); and 10 parameters  $w_1, \dots, w_{10}$ . So, each parameter  $w$  can be fine-tuned to match exactly each point (including noise).



# What has happened when $m = 10$ ?

But for points in the middle, the polynomial oscillates a lot!

By fitting each data point, values of  $w$  become too large!



This behaviour is known as **over-fitting**.





# The goal is to generalise, not overfit

Week 7:  
Regression 1

Dr Giuseppe  
Brandi


## Tip

The **goal** of machine learning is **good generalisation**; that is, the ability to predict new data.

Throwing more data to the problem helps, but...

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



We can say: “the size of the dataset should be  $5\times$  or  $10\times$  the size of the model.” But this is not good...

# Regularisation

We want the size of the model to be related to the complexity of the problem solved, not the size of the dataset, right?

**Regularisation is a technique to address over-fitting.**

It introduces a penalty term to the error function so that values of  $w$  do not get too large:

$$E(w) = \frac{1}{2} \sum_{i=1}^n (y - \hat{y})^2 + \alpha \text{ penalty}$$

# Regularisation

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

## ■ Ridge regression

$$E(w) = \frac{1}{2} \sum_{i=1}^n (y - \hat{y})^2 + \frac{\alpha}{2} \sum_{j=1}^m w_j^2$$

## ■ Lasso regression

$$E(w) = \frac{1}{2} \sum_{i=1}^n (y - \hat{y})^2 + \frac{\alpha}{2} \sum_{j=1}^m |w_j|$$

Parameter  $\alpha$  controls the **strength of the penalty**. Also, this technique is known as *weight decay*.

# Ridge regression

Week 7:  
Regression 1

Dr Giuseppe  
Brandi

## Solving it in Python


```
1 from sklearn.linear_model import Ridge
2 # Back to 10 input data points
3 x = np.linspace(0, 1, 10)
4 y = np.sin(2 * np.pi * x) +
5     rng.normal(scale=0.3, size=10)
6 # Create a ridge regression model with 10 params
7 model = make_pipeline(PolynomialFeatures(10),
8                       Ridge(alpha=0.0001))
9 model.fit(x[:, np.newaxis], y)
10 # Predict y-value for 1,000 data points
11 u = np.linspace(0, 1, 1000)
12 v = model.predict(u[:, np.newaxis])
```



# Ridge regression

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# Lasso regression


## Solving it in Python

```
1 from sklearn.linear_model import Lasso
2
3 # Create a lasso regression model with 10 params
4 model = make_pipeline(PolynomialFeatures(10),
5                         Lasso(alpha=0.0001,
6                               max_iter=100000))
7 model.fit(x[:, np.newaxis], y)
8 # Predict y-value for 1,000 data points
9 u = np.linspace(0, 1, 1000)
10 v = model.predict(u[:, np.newaxis])
```

# Lasso regression

Week 7:  
Regression 1

Dr Giuseppe  
Brandi



# Why is regularisation important?

A regulariser is one of the methods to **shrink parameter values towards zero**.

Thus, it is an opportunity to derive **sparse models**.

## Tip

We do not have to worry (much) about the size of the model with respect to the input data.

The *effective number of parameters* (that is, the non-zero ones) **automatically adapts** to the problem.

# Hyper-parameters

Let's recap the number of **hyper-parameters** that we have to tune to solve a regression problem:

- The choice of a basis function e.g., polynomial
- The model size, i.e., the number of parameters,  $m$
- The parameter  $\alpha$  of the regulariser (set  $\alpha = 0$  to disable it)
- The number of iterations (depending on the solver)

# Summary

We learned the fundamental of linear regression models, plus:

- What is an over-parameterised model?
- What is the over-fitting problem?
- What is regularisation?