



Week 9: Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

Week 9: Classification

Classification Algorithms

Dr Giuseppe Brandi

Northeastern University London



What is classification?

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

A type of supervised learning, where each input example has a known target.

- Given an input vector x , assign to it one of K discrete classes (or labels) C_k , $k = 1, 2, \dots, K$
- The mapping is unique (one class only)

Geometrically, we can view classification as dividing the input space into *regions*.

Example problems

Example, real-world problems include:

| <i>Application</i> | <i>Input</i> | <i>Output</i> |
|---------------------|--------------|---------------------|
| Spam detection | E-mails | True or false |
| Text recognition | Images | Characters & digits |
| Diagnostics | Symptoms | Disease(s) |
| Fraudulent activity | Transactions | True or false |



One-hot encodings

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

It is convenient to encode each class (or label) as “1-of- K ”:

- Each label is represented as a vector of length K
- The i -th label is a vector of zeros with a 1 at position i



One-hot encodings

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

For example, if $C = \{\text{cat}, \text{dog}, \text{mouse}\}$, we can define:

- cat = (1, 0, 0),
- dog = (0, 1, 0), and
- mouse = (0, 0, 1)

Intuitively, we can interpret such vectors as the probability of an input, say x , being a cat, dog, or a mouse:

$$(p(\text{cat}|x), p(\text{dog}|x), p(\text{mouse}|x))$$

The Linear Probability Model (LPM)

The Linear Probability Model (LPM) is a simple model, which entails applying linear regression to binary outcomes:

- **LPM model:**

$$P(y = 1 | x) = w_0 + w_1x_1 + w_2x_2 + \cdots + w_kx_k$$

- Each coefficient w_i represents the change in the probability of $y = 1$ for a one-unit change in x_i , assuming a linear relationship.
- **Interpreting Coefficients:** If $w_1 = 0.2$, a one-unit increase in x_1 raises the probability of $y = 1$ by 20%.
- **Intercept:** Represents the base probability when all predictors are zero.

Advantages of Linear Probability Model (LPM)

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

- **Simplicity:** Uses ordinary least squares (OLS) linear regression, easy to implement and understand.
- **Interpretability:** Coefficients represent marginal effects, providing straightforward interpretation (e.g., a 0.2 coefficient implies a 20% increase in probability).
- **Computational Efficiency:** LPM can be calculated quickly, even for large datasets.



Disadvantages of Linear Probability Model (LPM)

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

- **Probability Bounds:** Predictions can fall outside the 0-1 range, producing probabilities below 0 or above 1.
- **Non-Linearity:** Many relationships in binary outcomes are nonlinear, making LPM less accurate.
- **Sensitivity to Outliers:** Outliers can disproportionately affect predictions due to linearity.

Logistic regression

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

We know from linear regression models that:

- $y(x, w)$ is a linear function of parameters w
- $y(x) = w_0 + w_1x_1 + w_2x_2 + \dots = w_0 + w^T x$

But $y(x)$ is a real number, and we want to predict discrete labels; or better, $p(C_k|x)$ that are between 0 and 1.

Logistic regression

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

But $y(x)$ is a real number, and we want to predict discrete labels; or better, $p(C_k|x)$ that are between 0 and 1.

So, we apply a non-linear function to squash y :

$$p(C_k|x) = y(x) = f(w_0 + \mathbf{w}^T \mathbf{x})$$

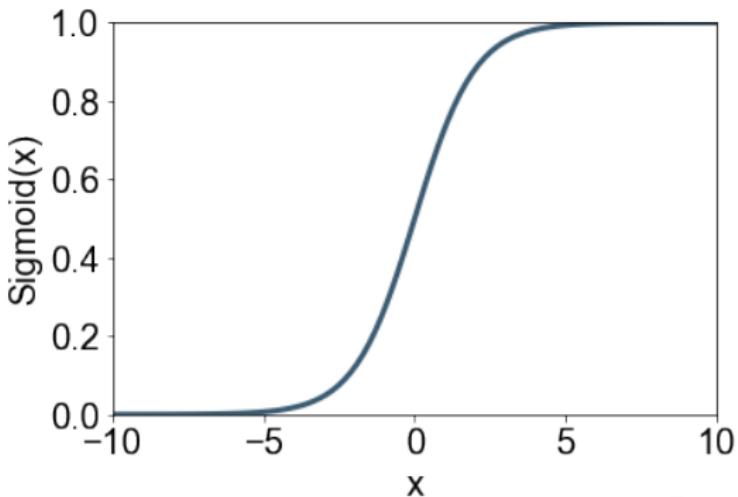
where f is:

- The sigmoid function for 2-class problems
- The softmax function for multi-class ones

Sigmoid function

The sigmoid function has a characteristic S-shape, given by:

$$f(x) = \frac{1}{1 + e^{-x}}$$



Logit Regression

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

- **Probability Constraints:** Ensures predictions remain between 0 and 1 using the logistic (sigmoid) function.
- **Nonlinear Relationship:** Better captures the underlying structure in binary outcomes, especially at the extremes (near 0 or 1).
- **More Robust:** Less sensitive to extreme values compared to LPM.

Functional Form:

$$P(y = 1|x) = \frac{1}{1 + e^{-(w_0 + w^T x)}}$$

Softmax Regression – Ideal for Multiclass Problems

- **Multiclass Capability:** Softmax is specifically designed for cases with multiple classes, generalizing Logit to multiple categories.
- **Probability Distribution:** Produces a probability for each class that sums to 1, ensuring interpretability and valid probability outputs.

Functional Form:

$$P(y = j|x) = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}, \quad j = 1, \dots, K$$

Example

Consider a 4-class classification problem given 300 training data points.

Generating sample data

```
1 from sklearn.datasets import make_blobs as blobs
2
3 X, y = blobs(n_samples=300,
4                 centers=4,
5                 n_features=2,
6                 random_state=0,
7                 cluster_std=1.0)
```

Example

Consider a 4-class classification problem given 300 training data points.

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

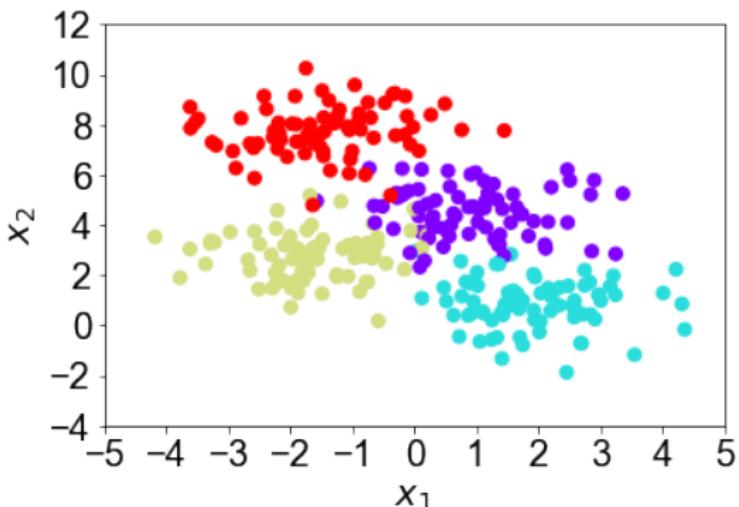
Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix



Solving it in Python

Consider a 4-class classification problem given 300 training data points.

Logistic regression

```
1 from sklearn.linear_model import LogisticRegression  
2  
3 # Create a logistic regression model  
4 model = LogisticRegression()  
5 # Fit the training data  
6 model.fit(X, y);
```

Solving it in Python

By calling `model.predict(samples)` on new data (`samples`), we can predict their label (in our case, colour).

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

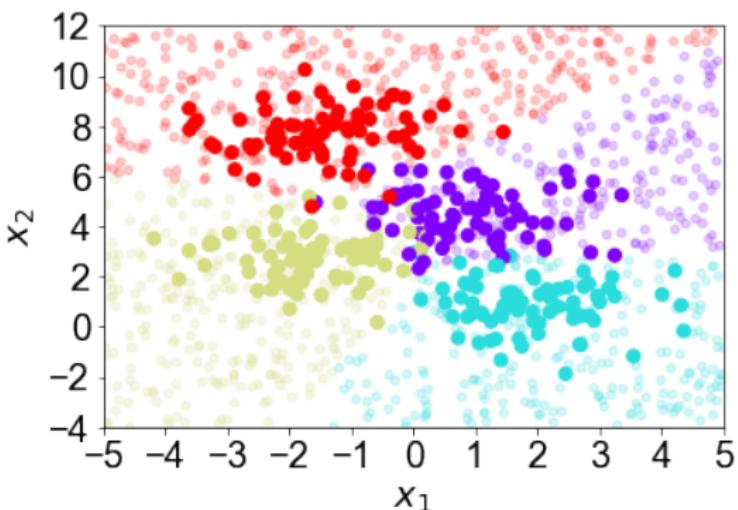
Linear
Probability
Model

Logit Model

Decision
Trees

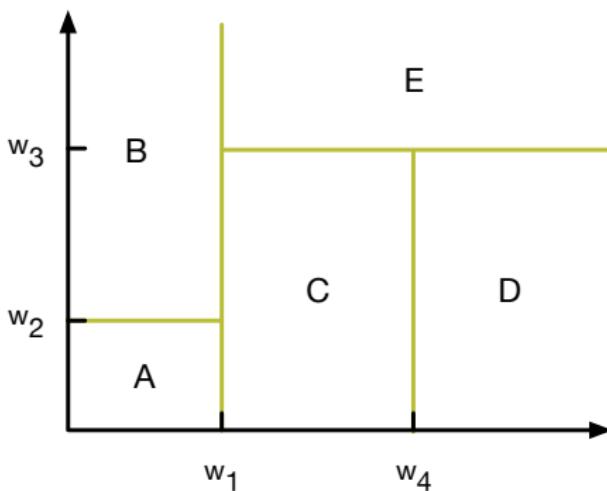
Performance
Metrics

Appendix



Decision trees

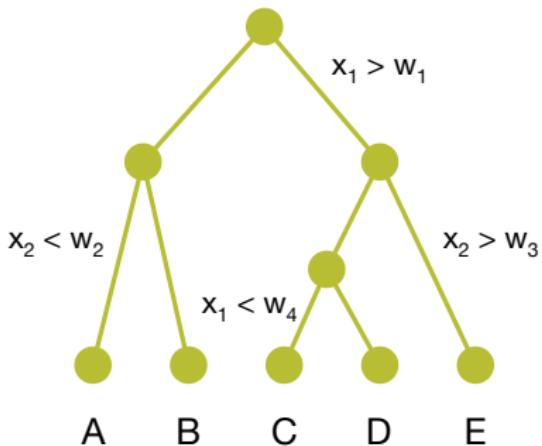
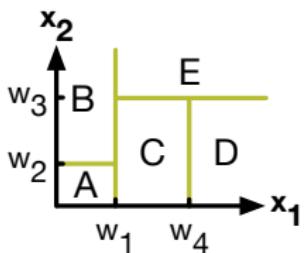
Decision trees partition the input space into cuboids (i.e., squares, cubes) and then assign a simple model to each region.¹



¹Example from C. Bishop's *Pattern recognition and machine learning*. ☺☺☺

Inference using decision trees

A decision tree makes decisions by following a path through questions, dividing data at each step. Each question splits the data into smaller parts until reaching a final decision.



Training decision trees

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

How to determine the best structure of a decision tree?

- Even with fixed number of nodes, it can be computationally infeasible – too many combinations
- Greedy construction starts at a root node and adds nodes one at a time.

Then, the question is:

- When to stop adding nodes? In other words,
- How deep should the decision tree be?

Solving it in Python

Consider a 4-class classification problem given 300 data points.

Logistic regression

```
1 from sklearn.tree import DecisionTreeClassifier  
2  
3 # Create a decision tree of depth 4  
4 classifier = DecisionTreeClassifier(max_depth=4)  
5 # Fit the training data  
6 classifier.fit(X, y);
```

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

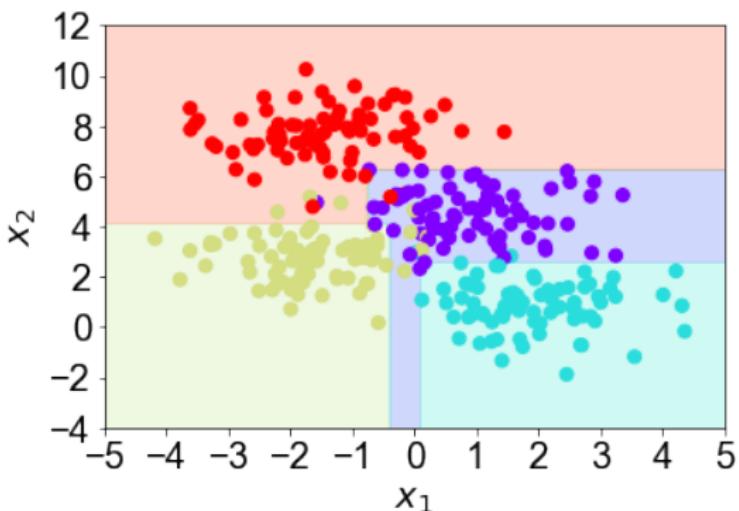
Decision
Trees

Performance
Metrics

Appendix

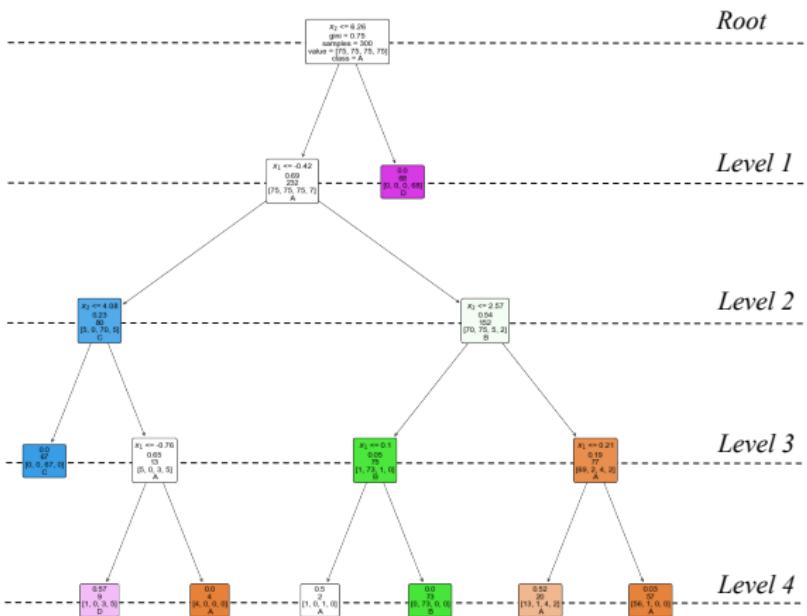
Solving it in Python

Let's visualise the regions found by the classifier:



Solving it in Python

Let's visualise the decision process:



Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

**Decision
Trees**

Performance
Metrics

Appendix



Improving Decision Trees: Pruning and Bagging

Challenges with Decision Trees

- Decision trees are prone to overfitting, especially when fully grown.
- They can have high variance, making predictions sensitive to training data fluctuations.

Two Key Techniques to Improve Decision Trees

- **Pruning:** Simplifies the tree by removing unnecessary branches, reducing overfitting and improving generalization.
- **Bagging:** Creates an ensemble of trees by training multiple trees on different subsets of data and averaging predictions, reducing variance.

Pruning in Decision Trees

What is Pruning?

- Pruning reduces the size of a decision tree by removing sections that do not provide significant predictive power.
- This process is typically done after the tree has been fully grown.

Why Use Pruning?

- **Reduce Overfitting:** By cutting back branches that capture noise, pruning helps the model generalize better to unseen data.
- **Improve Interpretability:** A pruned tree is simpler and easier to interpret.

Bagging

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

Bagging is an ensemble technique capable of reducing the variance of an estimator (e.g., a decision tree):

- A *committee* of classifiers, each fitting subsets of the data set
- Aggregating individual predictions (by voting or averaging) to make a final prediction.

Works well, but assumes individual models are uncorrelated

Solving it in Python

Consider a 4-class classification problem given 300 data points.

Logistic regression

```
1 from sklearn.ensemble import BaggingClassifier  
2  
3 tree = DecisionTreeClassifier()  
4 classifier = BaggingClassifier(tree,  
5                                 n_estimators=100,  
6                                 max_samples=0.8,  
7                                 random_state=1)  
8 # Fit the training data  
9 baggingclassifier.fit(X, y);
```

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

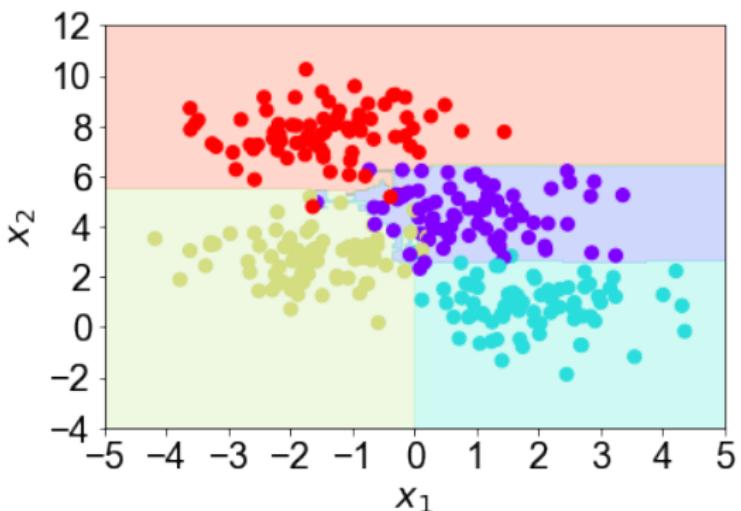
Decision
Trees

Performance
Metrics

Appendix

Solving it in Python

Let's visualise the regions found by bagging:



Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix



Limitations of decision trees

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

Decision trees are **easy to interpret**. However,

- High sensitivity to input data
- Regions are squares,
 - Boundaries are aligned to axes
 - Can you think of a data set where trees are not appropriate?

Performance Metrics for Classification

To evaluate and compare the effectiveness of classification models, we use several performance metrics:

- **Accuracy:** Measures the overall proportion of correct predictions among total predictions.
- **Precision:** Indicates the proportion of true positive predictions out of all positive predictions made by the model. Useful when the cost of false positives is high.
- **Recall (Sensitivity):** Measures the proportion of true positives out of all actual positives. Important when missing positive cases (false negatives) has a high cost.
- **F1 Score:** Combines precision and recall into a single metric by calculating their harmonic mean.

Performance Metrics for Classification

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

| | | true class |
|-----------------|-------|--|
| | | True False |
| predicted class | True | True Positives (TP) |
| | False | False Positives (FP) False Negatives (FN) |
| | | True Negatives (TN) |

$$PR = \frac{TP}{TP+FP}$$

$$RE = \frac{TP}{TP+FN}$$

$$CA = \frac{TP+TN}{TP+TN+FP+FN}$$

$$F_1 = \frac{2TP}{2TP+FP+FN}$$

Formulas for Performance Metrics: Recap

■ Accuracy:

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Predictions}}$$

■ Precision:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

■ Recall (Sensitivity):

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

■ F1 Score:

$$\text{F1 Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Appendix

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

Appendix: Naive Bayes

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

We want to find the conditional probability distribution:

$$p(C_k | \text{features})$$

One approach is to compute it using Bayes' theorem:

$$p(C_k | \text{features}) \propto p(\text{features} | C_k) p(C_k)$$

So, we need to compute:

- $p(C_k)$, otherwise known as the *prior* probability of C_k
- $p(\text{features} | C_k)$, the class-conditional probability of x

Why Naive Bayes?

The prior $p(C_k)$ can be estimated from the data set:

- How many instances of C_k exist in a dataset of size n .

But it is hard to estimate:

$$p(x_1, x_2, \dots, x_n | C_k)$$

Unless we simplify the problem by assuming independence:

$$p(x_1, x_2, \dots, x_n | C_k) = p(x_1 | C_k)p(x_2 | C_k) \cdots p(x_n | C_k)$$

Computing $p(x|C_k)$

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

We make assumptions about input features x_i , based on whether values are *discrete* or *continuous*.

If continuous, we typically assume a Gaussian (or normal) distribution. For each class C_k , $k = 1, 2, \dots, K$:

- Filter all x whose label is C_k
- For each x_i in x
 - Compute mean, μ_{x_i} , and standard deviation, σ_{x_i}
 - Then, $p(x_i|C_k) \sim N(\mu_{x_i}, \sigma_{x_i})$

Solving it in Python

Consider a 4-class classification problem given 300 training data points.

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix

Gaussian Naive Bayes

```
1 from sklearn.naive_bayes import GaussianNB  
2  
3 # Create a Gaussian Naive Bayes model  
4 model = GaussianNB()  
5 # Fit the training data  
6 model.fit(X, y);
```

Solving it in Python

By calling `model.predict(samples)` on new data (`samples`), we can predict their label (in our case, colour).

Week 9:
Classification

Dr Giuseppe
Brandi

Introduction

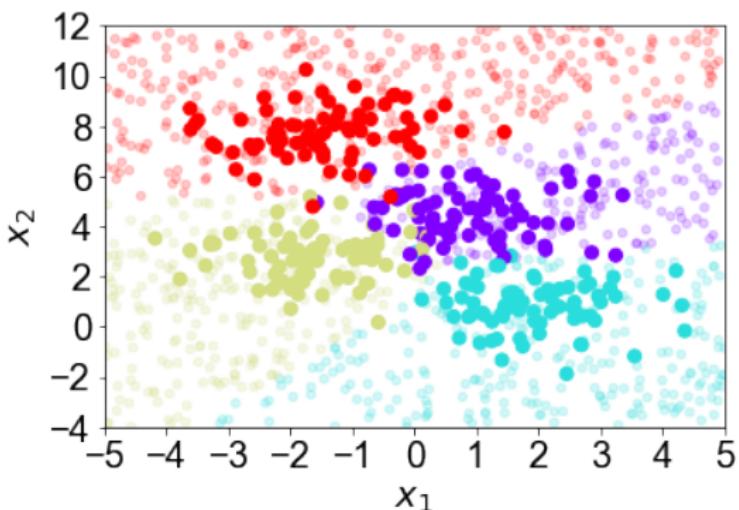
Linear
Probability
Model

Logit Model

Decision
Trees

Performance
Metrics

Appendix



Solving it in Python

Given $x = (0.74, 6.45)$, we computed:

- $p(C_1|x) = 0.72$
- $p(C_2|x) = p(C_3|x) = 0.00$
- $p(C_4|x) = 0.28$

