



Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Week 10: Classification 2

Dr Giuseppe Brandi

Northeastern University London



Outline

Week 10:
Classification
2

Dr Giuseppe
Brandi

- K-Nearest Neighbors (KNN) algorithm
- Classification reports
- Model selection & hyper-parameter tuning
- Feature selection & engineering

KNN

Example

Classification
report

Model
selection

Feature
Engineering



What is K-Nearest Neighbors (KNN)?

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

- K-Nearest Neighbors (KNN) is a simple, intuitive, and non-parametric machine learning algorithm.
- It can be used for both **classification** and **regression** tasks.
- KNN classifies or predicts based on the "**closeness**" or **similarity** of data points, measured by a distance metric (e.g., Euclidean distance).



How KNN Works (Step-by-Step)

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

- 1 Choose the number of neighbors k .
- 2 Calculate the distance between the target point and all points in the training set.
- 3 Select the k closest points (neighbors).
- 4 For classification: Assign the class most common among neighbors.
For regression: Take the average value of the neighbors.



Choosing the Value of k

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering



Distance Metrics in KNN

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

- Different metrics can be used to calculate the "distance" or "similarity" between points.
- Common distance metrics include:
 - **Euclidean Distance:** For continuous data, calculates straight-line distance.
 - **Manhattan Distance:** For categorical or grid-like data, calculates the sum of absolute differences.
 - **Cosine Similarity:** For text and high-dimensional data, measures angle between vectors.

Advantages and Disadvantages of KNN

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

■ Advantages

- **Simple to understand** and easy to implement.
- **No training phase:** All computation happens at prediction time.
- **Flexible with distance metrics**, making it adaptable to various data types.

■ Disadvantages

- **Computationally expensive** for large datasets, as it requires calculating the distance to all points.
- **Sensitive to irrelevant or noisy features** and scales of the data.
- **Choice of k and distance metric** can significantly impact performance.

Applications of KNN

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

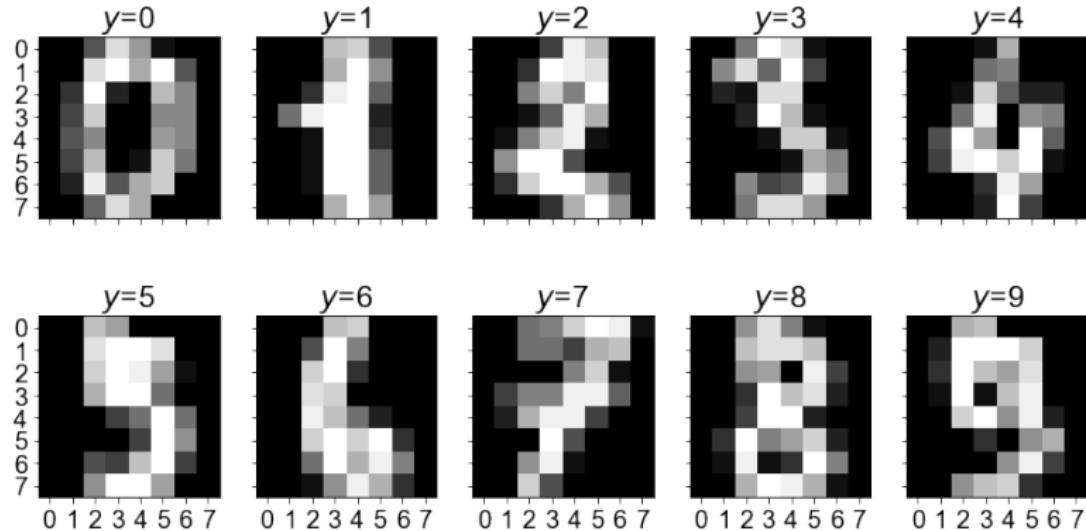
Model
selection

Feature
Engineering

- **Recommendation Systems:** Finding similar items or users.
- **Medical Diagnosis:** Predicting disease based on patient symptoms and historical data.
- **Image and Text Classification:** Classifying objects in images or documents based on similarity.
- **Anomaly Detection:** Identifying outliers in data, such as fraud detection.

Example dataset

1,797 8×8 images of digits (0-9):





The goal is to generalise, not overfit

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Tip

The **goal** of machine learning is **good generalisation**; that is, the ability to predict new data.



Splitting the data set in Python

Use 75% of data for *training* and 25% for *testing*; and **shuffle!**

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Splitting a dataset

```
1 from sklearn.model_selection
2     import train_test_split as split
3
4 train.X, test.X, train.y, test.y =
5             split(images,
6                     labels,
7                     test_size=0.25,
8                     random_state=123456789)
```

Training a classifier

As an example, let's use a ***k*-nearest neighbours** classifier:

Training a *k*-nearest neighbours classifier

```
1 from sklearn.neighbors import KNeighborsClassifier  
2  
3 # Create the classifier  
4 classifier = KNeighborsClassifier()  
5  
6 # Fit the training data  
7 classifier.fit(X=train.X, y=train.y);
```

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Test vs training accuracy

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

```
1 accuracy = classifier.score(test.X, test.y)
2 accuracy = accuracy * 100.0
3 print('Test accuracy is {:.2f}%'.format(accuracy))
```

Test accuracy is 97.33%

This is to be **contrasted** with a training accuracy of 99.33%.

Confusion matrix

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

A confusion matrix C is such that $C(i,j)$ equals the number of observations known to have label i and assigned label j .

- Correct predictions are on the **diagonal**
- Non-zeros not on the diagonal are *incorrect* predictions
- Each **row** represents a class (in our case, 0-9)

Confusion matrix in Python

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

```
1 from sklearn.metrics import confusion_matrix
2
3 # test.y are the known labels for each test image.
4 # But what does our model predict?
5 predictions = classifier.predict(test.X)
6
7 C = confusion_matrix(test.y, predictions)
```

Let's plot the confusion matrix.



Visualising a confusion matrix

Week 10:
Classification
2

Dr Giuseppe
Brandi

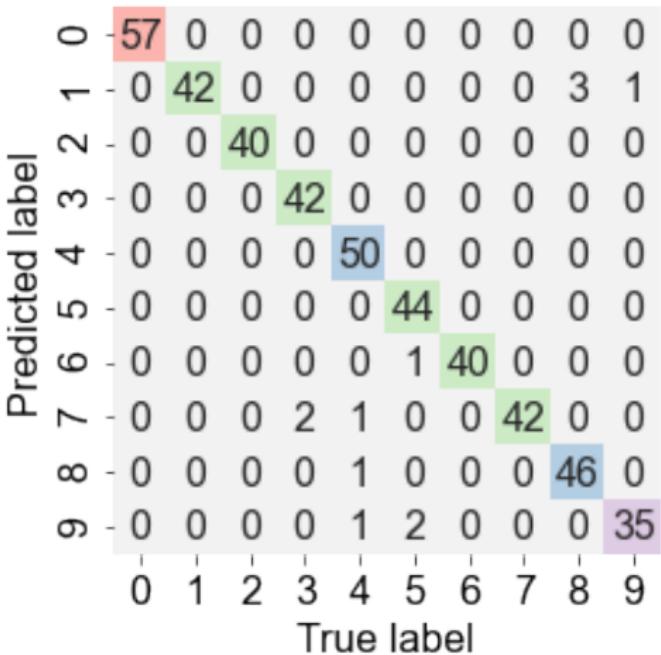
KNN

Example

Classification
report

Model
selection

Feature
Engineering



Classification report

Generating a classification report

```
1 from sklearn.metrics import classification_report
2 report = classification_report(test.y, predictions)
3 print(report)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	57
1	0.91	1.00	0.95	42
2	1.00	1.00	1.00	40
3	1.00	0.95	0.98	44
4	1.00	0.94	0.97	53
...				



Understanding a classification report

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

- **Precision** measures the fraction of correct predictions, incl. errors, for a label
- **Recall** measures a classifier's ability to predict correctly all instances of a label
- For example, for digit “1”
 - Precision is 91% – not all predictions were correct
 - Recall is 100% – all “1”s were classified correctly
 - So, our classifier mistaken other digits for a “1”

Model selection problem

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

In practical machine learning applications, we need to set a number of *hyper-parameters*

Besides tuning a particular model, we also want to consider a range of different types of model to find the best one for a particular application

Model selection problem

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

The typical approach is to split our data into three sets:

- Training set
- Validation set
- Test set

But what if the supply of data is limited? One solution is to use **cross-validation**

Cross validation

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

The main idea is to use **all data** for training and testing by splitting the data set into **K equally-sized folds** (i.e., data partitions)

Suppose that $K = 4$, resulting in 4 folds, S_1, S_2, S_3 and S_4

Then, for we perform $K = 4$ experiments:

- Train the model with S_2, S_3, S_4 and test with S_1
- Train the model with S_1, S_3, S_4 and test with S_2
- Train the model with S_1, S_2, S_4 and test with S_3
- Train the model with S_1, S_2, S_3 and test with S_4

and average the classification scores from each run.

Cross validation in Python

Creating a K-fold

```
1 from sklearn.model_selection import KFold,  
2                               cross_val_score  
3  
4 folds = KFold(n_splits=10, shuffle=True)  
5 classifier = KNeighborsClassifier()  
6 scores = cross_val_score(estimator=classifier,  
7                           X=images, y=labels,  
8                           cv=folds)  
9 print(f'{len(scores)} scores: {scores.mean():.2%}'  
10      +- {scores.std():.2%}' )
```

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Model comparison

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

It is difficult to know in advance which machine learning models will perform best for a given dataset.

This is particularly true when using a library like sklearn, where implementation details are hidden from us.

Let's compare the k -nearest neighbour classifier with another method:

- `DecisionTreeClassifier()`

Model comparison

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

```
1 estimators = {  
2     'K-nearest neighbours': KNeighborsClassifier(),  
3     'Decision Tree': DecisionTreeClassifier()  
4 }  
5  
6 folds = KFold(n_splits=10,  
7                 random_state=123456789,  
8                 shuffle=True)
```

Model comparison

Comparing three classifiers

```
1 for name, classifier in estimators.items():
2     # Cross-validate classifier and report result
3     scores = cross_val_score(estimator=classifier,
4                               X=images,
5                               y=labels,
6                               cv=folds)
7     print(f'{name:20s}: {scores.mean():.2%}
8           +- {scores.std():.2%}')
```

K-nearest neighbours: 98.72% +- 0.79%
Decision Tree : 86.09% +- 2.37%



Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Hyper-parameter tuning

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Each model has a number of (hyper-)parameters that can be tuned for a particular application and data set.

Examples include:

- Leaf size & num. neighbours for `KNeighborsClassifier()`
- Max tree depth for `DecisionTreeClassifier()`

Grid search

It is difficult to find the values of the important parameters of a model that provide the best generalisation performance.

Grid search is a method to try all possible combinations of the parameters of interest.

- Use the `sklearn.model_selection.GridSearchCV` method
- Uses K -fold cross-validation behind the scenes for finding the best combination



Grid search

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

```
1 from sklearn.model_selection import GridSearchCV
2 # Parameters to search, and possible values
3 grid = {
4     'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
5     'leaf_size':    [10, 20, 30, 40, 50]
6 }
7 g = GridSearchCV(classifier, grid, cv=10)
8 g.fit(X=train.X, y=train.y)
9 g.best_params_
```

```
{'leaf_size': 10, 'n_neighbors': 5}
```



Feature selection

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

We can use a **univariate feature selection** method:

- Checks for a *statistically significant relationship* between each feature and target
- Each feature is individually considered
- Features that are related with the highest confidence are selected



Feature selection in Python

Let's try to select the 42 most important features (out of 64) in our digits dataset

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Select 42 out of 64 pixels per image

```
1 # Define a two-part selection method
2 extractor = make_pipeline(
3     VarianceThreshold(threshold=0),
4     SelectKBest(k=42))
5
6 extractor.fit(train.X, train.y)
7 extracted = extractor.transform(train.X)
```

Classifier re-trained with extracted data has same accuracy!



Feature selection – outlook

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

Besides univariate feature selection, we can use:

- Model-based feature selection
 - Train another model, this time for feature selection
- Iterative feature selection
 - Try all combinations

Feature engineering

Consider the following data set. Feature *Group* cannot be used as an input to a machine learning algorithm directly.

Name	Group	Grade
Alex	A	70.1
Jane	A	95.0
John	B	65.8
Lucy	C	100.0

But we can engineer it.

Feature engineering

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

```
1 from sklearn.preprocessing import OneHotEncoder
2
3 # Select column 'group'
4 groups = df.group.values.reshape(-1, 1)
5 # One-hot encoding
6 encoder = OneHotEncoder(sparse=False)
7 g = encoder.fit_transform(groups)
```

```
1 0 0  # was A
1 0 0  # was A
0 1 0  # was B
0 0 1  # was C
```



Feature engineering

Similar to one-hot encodings, we can **discretise continuous variables** by splitting them into a fixed number of bins

Creating a one-hot encoding of a column

```
1 from sklearn.preprocessing import KBinsDiscretizer
2
3 binning = KBinsDiscretizer(n_bins=5,
4                             strategy='uniform',
5                             encode='onehot-dense')
```

```
1 0 0 0 0 # was 70.1
...
...
```





Summary

Week 10:
Classification
2

Dr Giuseppe
Brandi

KNN

Example

Classification
report

Model
selection

Feature
Engineering

- Classification reports – test accuracy & confusion matrix
- Model selection & hyper-parameter tuning
 - Cross-validation
- Feature selection & engineering